

Categorical variables

(Chriss Jordan Oboa)

Contents

Setup and packages	1
Loading in data	1
Categorical variables	1
How many cases are in this dataset? How many variables?	2
Visualization for categorical variables	2
Two categorical variables	4
Summary statistics for categorical variables	7
Two categorical variables	8

Setup and packages

For this activity, we will just need to make use of the `tidyverse` package. Load it below.

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --
## v ggplot2 3.3.5      v purrr   0.3.4
## v tibble  3.1.6      v dplyr   1.0.7
## v tidyr   1.1.4      v stringr 1.4.0
## v readr   2.1.1      v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

Loading in data

In the code chunk below, read in the dataset that is stored in the file `GSS_clean.csv`, which is located inside of the `data` folder here in RStudio Cloud. Give the dataset the name `GSS` and use the `read_csv` command.

```
GSS<-read.csv("data/GSS_clean.csv")
```

Categorical variables

This dataset is from the General Social Survey (GSS). From the GSS website,

“The General Social Survey (GSS) is a nationally representative survey of adults in the United States conducted since 1972. The GSS collects data on contemporary American society in order to monitor and explain trends in opinions, attitudes and behaviors. The GSS has adapted questions from earlier surveys, thereby allowing researchers to conduct comparisons for up to 80 years.”

You can find more information about the GSS here: <https://gss.norc.oregonstate.edu/About-The-GSS>.

Use `glimpse`, `head`, or `View` to take a look at the dataset.

How many cases are in this dataset? How many variables?

There are 2348 cases and 45 variables.

Name some variables in this dataset that are categorical? Govt_or_Private_employee Marital_status College_major

Visualization for categorical variables

There are two main ways to visualize one categorical variable: **bar charts**, and pie charts. We're going to focus on bar charts here, because they are easier to create in R. Pie charts are generally not recommended – see the discussion on pie charts in our textbook.


The variable we will focus on first is `marital_status` which gives the marital status of each respondent. To see the different categories or “levels” associated with this variable, we can do the following:

```
GSS %>%  
  distinct(marital_status)
```

```
##   marital_status  
## 1  Never married  
## 2      Separated  
## 3      Married  
## 4      Divorced  
## 5      Widowed  
## 6             <NA>
```


Next, we will again use `ggplot` to help us graph the distribution of a categorical variable in a bar plot.

```
ggplot(data = GSS, mapping = aes(x = marital_status)) +  
  geom_bar()
```



Below, try making your own bar chart of the `labor_force_status`.


```
ggplot(data = GSS, mapping = aes(y = labor_force_status)) +  
  geom_bar(width=0.5, fill="red")
```



Two categorical variables


If we want, we can make a barchart showing **two** categorical variables. To do this, we add another argument to our `aes()` function. We will tell R to fill in the `marital_status` bars with color, according to the `born_in_us` variable. This makes a segmented or stacked barchart.

```
ggplot(data = GSS, mapping = aes(x = marital_status, fill = born_in_us)) +  
  geom_bar()
```



We can also make this barchart a side-by-side or “dodged” barchart by providing the `position` argument to the `geom_bar` function.


```
ggplot(data = GSS, mapping = aes(x = marital_status, fill = born_in_us)) +  
  geom_bar(position = "dodge")
```



Try making a barchart that shows both the `govt_or_private_employee` and `self_emp_or_works_for_somebody` variables.

```
ggplot(data = GSS, mapping = aes(x = govt_or_private_employee, fill = self_emp_or_works_for_somebody)) +
  geom_bar(position = "dodge")
```

Warning: Ignoring unknown parameters: position



Summary statistics for categorical variables

The other part of Exploratory Data Analysis (EDA) is making summary statistics. For a categorical variable, this means making frequency tables and relative frequency tables. Let's explore how to do that.

We will use the pipe, `%>%`, to string together three main R functions:

- `summarize()`, which summarizes data (going from many rows down to just one or several)
- `mutate()`, which makes new variables in our dataset (sticking on one more column)
- `group_by()`, which works especially well with `summarize()` to group data for operations

Let's start with a short data pipeline.

```
GSS %>%
  summarize(n = n())
```

```
##      n
## 1 2348
```

What does this code tell us? how many variables are in the dataset.

We can add another piece to the pipeline to make it more interesting.

```
GSS %>%
  group_by(marital_status) %>%
  summarize(n = n())
```

```
## # A tibble: 6 x 2
##   marital_status      n
##   <chr>           <int>
## 1 Divorced         403
```

```
## 2 Married          998
## 3 Never married    670
## 4 Separated        75
## 5 Widowed          200
## 6 <NA>              2
```

```
#frequency table
```

What does this code give us? It gives us all type of marital status and the number of people that fall into each category.

Finally, we might want to add one more piece to the pipeline to create a relative frequency table,

```
GSS %>%
  group_by(marital_status)%>%
  summarize(n = n()) %>%
  mutate(proportion = n/sum(n))
```

```
## # A tibble: 6 x 3
##   marital_status      n proportion
##   <chr>          <int>     <dbl>
## 1 Divorced        403    0.172
## 2 Married         998    0.425
## 3 Never married   670    0.285
## 4 Separated       75    0.0319
## 5 Widowed         200    0.0852
## 6 <NA>            2    0.000852
```

Two categorical variables

We can extend this to two categorical variables by adding a second variable into our `group_by()` code. Let's make a relative frequency table showing the relationship between `marital_status` and `general_happiness`.

```
GSS %>%
  group_by(marital_status, general_happiness)%>%
  summarize(n = n()) %>%
  mutate(proportion = n/sum(n))
```

`summarise()` has grouped output by 'marital_status'. You can override using the `.groups` argument.

```
## # A tibble: 20 x 4
## # Groups:   marital_status [6]
##   marital_status general_happiness      n proportion
##   <chr>          <chr>          <int>     <dbl>
## 1 Divorced      Not too happy      84    0.208
## 2 Divorced      Pretty happy     242    0.600
## 3 Divorced      Very happy       77    0.191
## 4 Married       Not too happy     61    0.0611
## 5 Married       Pretty happy     504    0.505
## 6 Married       Very happy      432    0.433
## 7 Married       <NA>              1    0.00100
## 8 Never married Not too happy    135    0.201
## 9 Never married Pretty happy    409    0.610
## 10 Never married Very happy     124    0.185
## 11 Never married <NA>              2    0.00299
## 12 Separated    Not too happy     19    0.253
## 13 Separated    Pretty happy     40    0.533
## 14 Separated    Very happy      15    0.2
```



```
## 15 Separated      <NA>                1    0.0133
## 16 Widowed       Not too happy        37    0.185
## 17 Widowed       Pretty happy        111   0.555
## 18 Widowed       Very happy          52    0.26
## 19 <NA>          Pretty happy          1    0.5
## 20 <NA>          Very happy           1    0.5
```

What happens if you switch the order of the variables in the `group_by()`? of all the people that indicate they are divorced, only 20% are happy.

```
GSS %>%
  group_by(general_happiness, marital_status)%>%
  summarize(n = n()) %>%
  mutate(proportion = n/sum(n))
```

``summarise()`` has grouped output by 'general_happiness'. You can override using the ``.groups`` argument

```
## # A tibble: 20 x 4
## # Groups:   general_happiness [4]
##   general_happiness marital_status      n proportion
##   <chr>             <chr>         <int>     <dbl>
## 1 Not too happy    Divorced          84    0.25
## 2 Not too happy    Married           61    0.182
## 3 Not too happy    Never married    135    0.402
## 4 Not too happy    Separated        19    0.0565
## 5 Not too happy    Widowed          37    0.110
## 6 Pretty happy     Divorced        242    0.185
## 7 Pretty happy     Married         504    0.386
## 8 Pretty happy     Never married    409    0.313
## 9 Pretty happy     Separated        40    0.0306
## 10 Pretty happy    Widowed         111    0.0849
## 11 Pretty happy    <NA>             1    0.000765
## 12 Very happy     Divorced         77    0.110
## 13 Very happy     Married          432    0.616
## 14 Very happy     Never married    124    0.177
## 15 Very happy     Separated        15    0.0214
## 16 Very happy     Widowed          52    0.0742
## 17 Very happy     <NA>             1    0.00143
## 18 <NA>           Married           1    0.25
## 19 <NA>           Never married     2    0.5
## 20 <NA>           Separated         1    0.25
```

of all the people that indicate they are divorced, only 20% are happy.